VIT®
Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)

VIT-SanDisk Hackathon 2026 - Project Idea Submission

1. Project Overview

Track Name: Track 1: AI / ML

Theme Name: Intelligent Storage Orchestration & Predictive Maintenance

Project Name: GuardianDrive

Tagline: AI-Powered Predictive Health, Risk-Aware Tiering, and Multi-Cloud Orchestration

Live Project URL: https://divyanshupatel.com/GuardianDrive-sandisk/

2. Problem Statement

Modern storage environments in both consumer and enterprise sectors face four critical, interconnected challenges that result in data loss, inflated costs, and inefficiency.

A. Silent Drive Failures & Reactive Maintenance

- The Issue: Storage drives (HDDs/SSDs) often degrade silently. Traditional S.M.A.R.T. metrics are cryptic and ignored by 87% of users until catastrophic failure occurs.

- The Impact: Data loss is reactive; backups are often triggered *after* corruption has started or are not comprehensive enough for the specific data risks.

B. The Cloud Cost Spiral

- The Issue: Enterprises and power users waste 40-60% of cloud storage budgets on improperly tiered data.

  - Hot data (frequently accessed) is often trapped in slow, cold archives.

  - Cold data (rarely accessed logs/backups) sits in expensive premium SSD tiers.

- The Impact: Massive, unnecessary monthly bills due to a lack of granular, file-level visibility.

C. Fragmented Storage Landscape

- The Issue: Infrastructure is siloed between Local NVMe/HDDs, External Arrays, and multiple Cloud Providers (AWS, Azure, GCP).

- The Impact: There is no "single pane of glass" to orchestrate data movement based on health, cost, and performance across these disparate layers.

D. "Data Gravity" & Compression Guesswork

- The Issue: Moving large datasets is expensive (egress fees) and slow. Users blindly compress files without knowing if it will yield significant space savings.

- The Impact: Wasted CPU cycles compressing incompressible media (JPEGs, MP4s) and missed opportunities to compress text-heavy logs/databases.

GuardianDrive Solution:

An AI-native orchestration platform that unifies predictive maintenance, intelligent tiering, and multi-cloud arbitrage into a single, automated workflow.

3. Detailed Architecture / Block Diagram

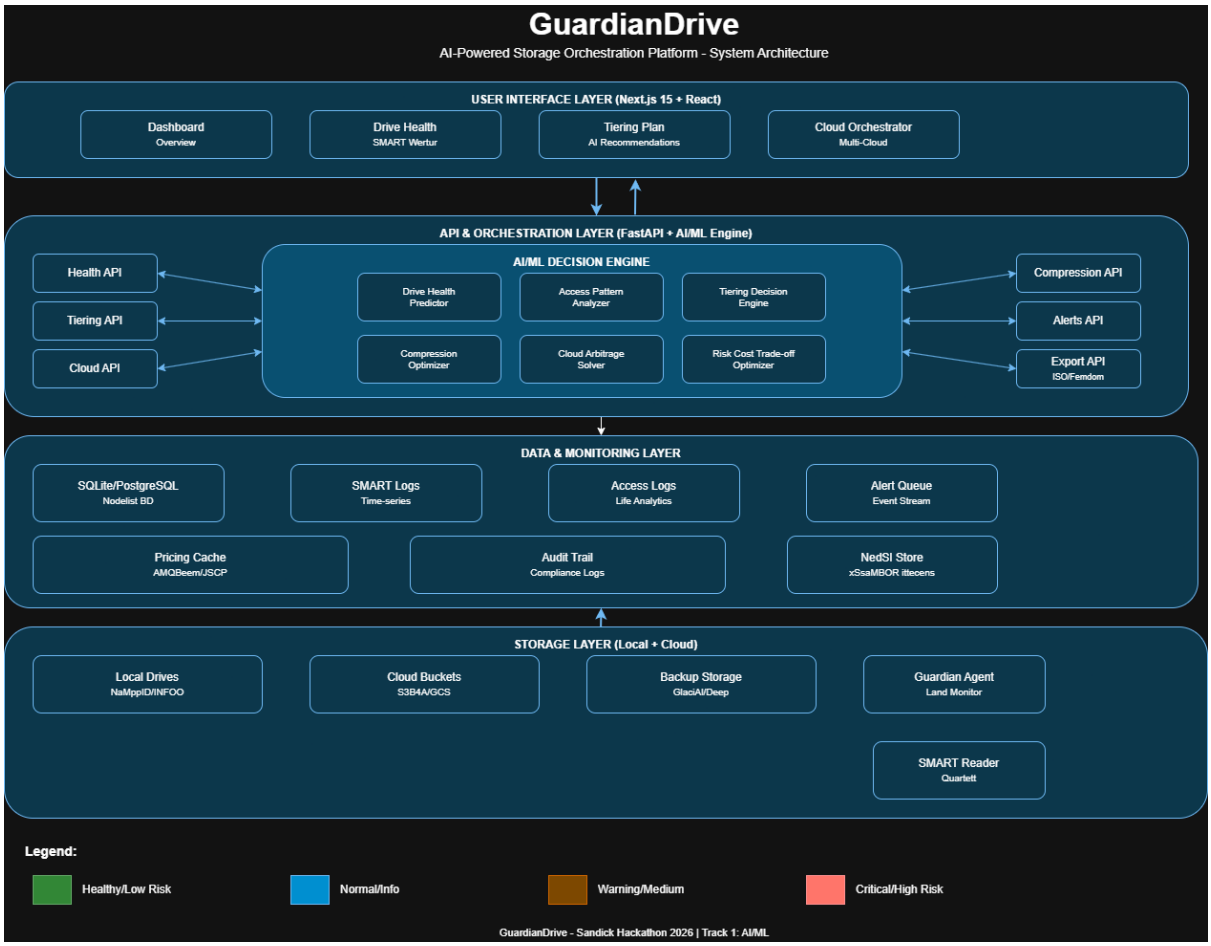

System Architecture Overview:

Layer 1: USER INTERFACE LAYER

- React 19 + Vite + Tailwind Obsidian Dashboard

- Components: Overview, Health Monitor, Smart Tiering, Compression, Cloud Ops

Layer 2: API & ORCHESTRATION LAYER

- FastAPI + Python 3.11 Backend

- Modules: Health API, Tiering API, Cloud API

- AI/ML Decision Engine (The "Brain"):

  1. Drive Health Model (XGBoost Classifier)

  2. Access Pattern Model (K-Means Clustering)

  3. Cost Solver (Weighted Scalarization)

Layer 3: DATA LAYER

- SQLite/PostgreSQL Metadata

- SMART Telemetry DB

- File Access Logs

- Cloud Pricing Cache

Layer 4: STORAGE LAYER

- Local NVMe/HDD

- External Arrays

- NAS / DAS

Layer 5: EXTERNAL SERVICES

- AWS S3 API

- Azure Blob API

- GCP Storage API

4. Detailed Methodology

Our approach combines supervised learning for health prediction, unsupervised learning for data classification, and mathematical optimization for cost placement.

Phase 1: Predictive Drive Health (The "Guardian")

- Data Source: Trained on the Backblaze Hard Drive Dataset (100k+ drives, 5 years of history).

- Feature Engineering: Uses raw SMART attributes (Reallocated Sectors, Seek Error Rate, Power-On Hours, Temperature) and time-series features (Rolling averages) to detect degradation trends.

- Model: XGBoost Classifier.

- Objective: Predict failure within a 14-day horizon.

- Output: A granular Health Score (0-100) and a Risk Class (Low/Medium/Critical).

- Innovation: Unlike simple "Pass/Fail" SMART flags, our model provides a probabilistic "days to failure" countdown.

Phase 2: Intelligent Data Classification (The "Auditor")

- Problem: Files are not equal. A 10GB log file is different from a 10GB 4K video.

- Method: K-Means Clustering.

- Features: Recency (days since last access), Frequency (access count), and Size.

- Clusters: The model automatically segments files into 4 tiers:

1. HOT: Frequently accessed, recent. (Keep on NVMe)
2. WARM: Occasional access. (Move to HDD/SATA)
3. COLD: Rare access. (Move to Cloud Cool tier)
4. ARCHIVE: Never accessed, compliance data. (Move to Glacier/Deep Archive)

Phase 3: Risk-Cost Optimizer (The "Broker")

- Problem: Minimizing cost often increases risk. We need a balance.

- Method: Weighted Scalarization Optimization.

- Formula: Score = w1 * Norm(Cost) + w2 * Norm(Risk) + w3 * Norm(Latency)

- Execution: The system generates 3 distinct strategies for the user:

1. Conservative: Max redundancy, High Performance (High Cost)
2. Balanced: (Recommended) Optimal trade-off (Medium Cost)
3. Aggressive: Lowest possible cost, Cloud Archive focus (Low Cost)

Phase 4: Compression ROI Analysis

- Logic: Before compressing, we calculate the Return on Investment (ROI).

- Calculation: If (Predicted Savings > Compute Cost), compression is recommended.

- Tech: Uses Zstandard (zstd) for variable compression levels based on the file type.

Phase 5: Multi-Cloud Arbitrage

- Method: Real-time query of AWS/Azure/GCP pricing APIs.

- Action: Compares current storage costs vs. competitors. Auto-generates Terraform scripts to migrate buckets if a cheaper provider is found. Generates S3 Lifecycle Policies to automate transitions.

5. Deliverables Expected & Achieved

A. Functional Web Dashboard (MVP)

- A fully responsive, dark-mode dashboard built with React & Vite.

- Includes real-time storage health monitoring and visual drive life expectancy gauges.

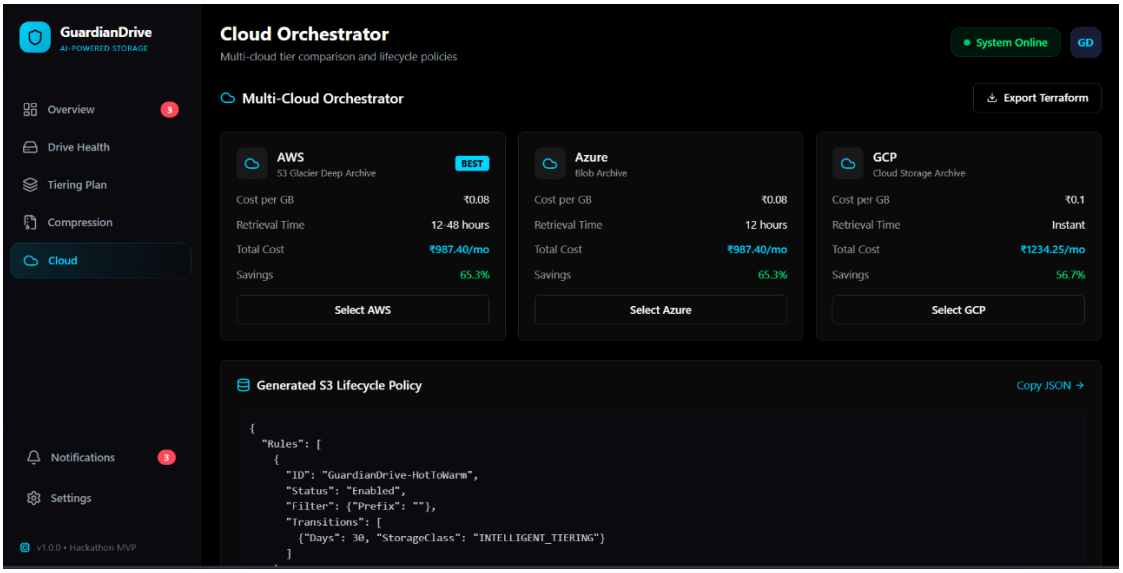

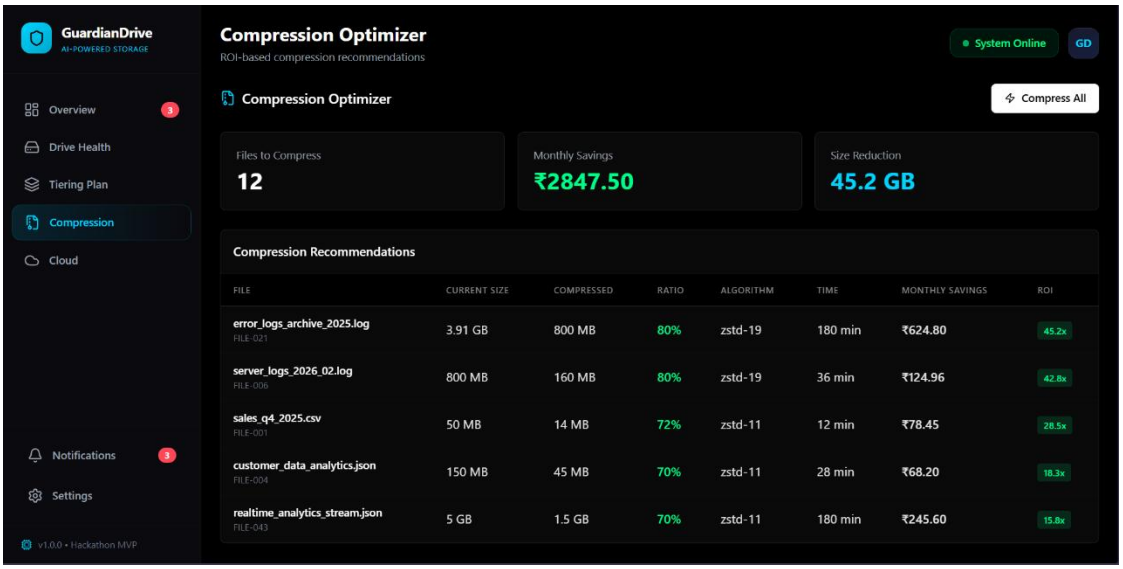

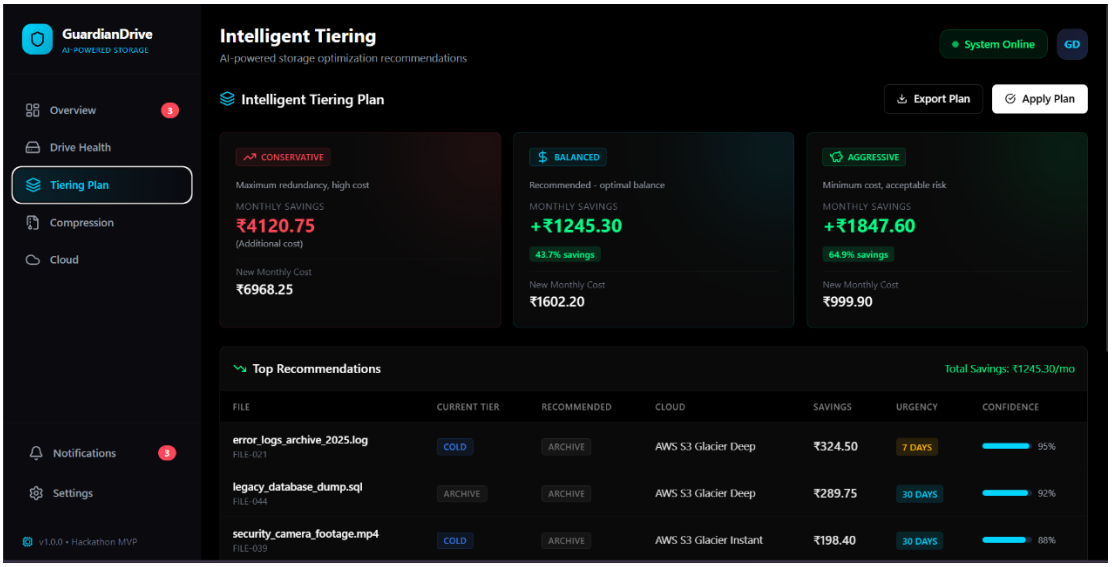
GuardianDrive
AI-POWERED STORAGE
Overview 3
Drive Health
Tiering Plan
Compression
Cloud
Notifications 3
Settings
v1.0.0 • Hackathon MVP
Intelligent Tiering
AI-powered storage optimization recommendations
System Online
GD
Intelligent Tiering Plan
Export Plan
Apply Plan
CONSERVATIVE
Maximum redundancy, high cost
MONTHLY SAVINGS
₹4120.75
(Additional cost)
New Monthly Cost
₹6968.25
BALANCED
Recommended - optimal balance
MONTHLY SAVINGS
+₹1245.30
43.7% savings
New Monthly Cost
₹1602.20
AGGRESSIVE
Minimum cost, acceptable risk
MONTHLY SAVINGS
+₹1847.60
64.9% savings
New Monthly Cost
₹999.90
Top Recommendations
Total Savings: ₹1245.30/mo
FILE
CURRENT TIER
RECOMMENDED
CLOUD
SAVINGS
URGENCY
CONFIDENCE
error_logs_archive_2025.log
FILE-021
COLD
ARCHIVE
AWS S3 Glacier Deep
₹324.50
7 DAYS
95%
legacy_database_dump.sql
FILE-044
ARCHIVE
ARCHIVE
AWS S3 Glacier Deep
₹289.75
30 DAYS
92%
security_camera_footage.mp4
FILE-039
COLD
ARCHIVE
AWS S3 Glacier Instant
₹198.40
30 DAYS
88%

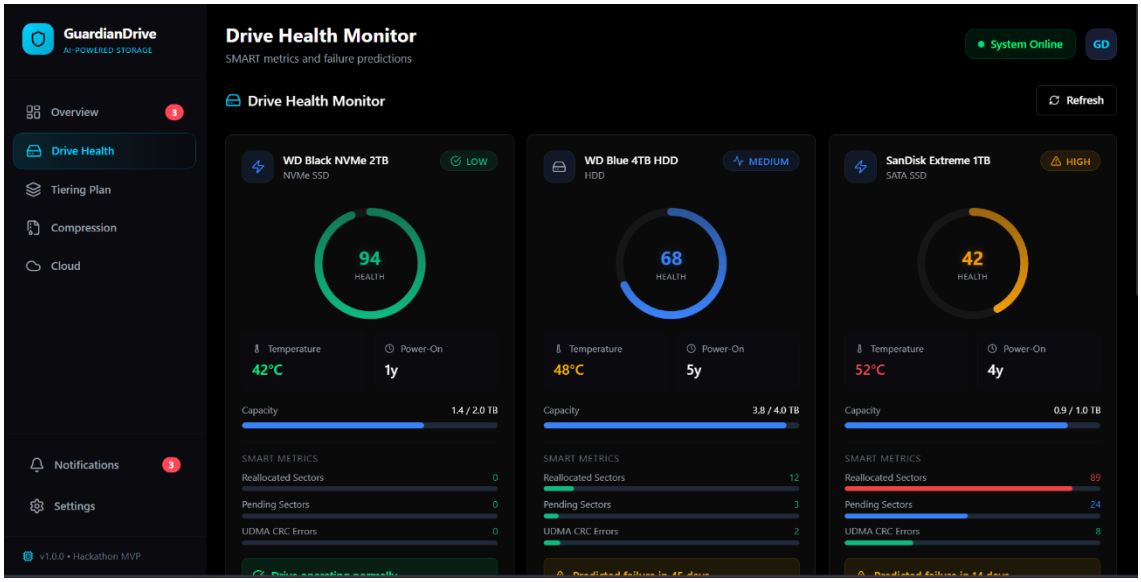
GuardianDrive
AI-POWERED STORAGE
Overview 3
Drive Health
Tiering Plan
Compression
Cloud
Notifications 3
Settings
v1.0.0 • Hackathon MVP
Drive Health Monitor
SMART metrics and failure predictions
System Online
GD
Drive Health Monitor
Refresh
WD Black NVMe 2TB
NVMe SSD
LOW
94
HEALTH
Temperature
42°C
Power-On
1y
Capacity
1.4 / 2.0 TB
SMART METRICS
Reallocated Sectors 0
Pending Sectors 0
UDMA CRC Errors 0
WD Blue 4TB HDD
HDD
MEDIUM
68
HEALTH
Temperature
48°C
Power-On
5y
Capacity
3.8 / 4.0 TB
SMART METRICS
Reallocated Sectors 12
Pending Sectors 3
UDMA CRC Errors 2
SanDisk Extreme 1TB
SATA SSD
HIGH
42
HEALTH
Temperature
52°C
Power-On
4y
Capacity
0.9 / 1.0 TB
SMART METRICS
Reallocated Sectors 89
Pending Sectors 24
UDMA CRC Errors 8

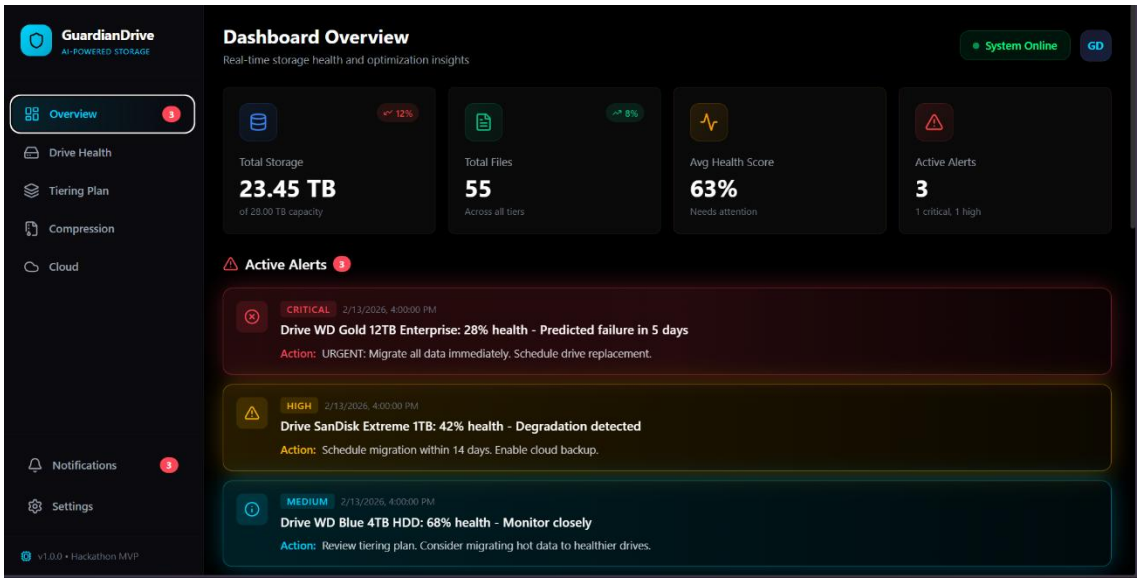
GuardianDrive
AI-POWERED STORAGE
Overview 3
Drive Health
Tiering Plan
Compression
Cloud
Notifications 3
Settings
v1.0.0 • Hackathon MVP
Dashboard Overview
Real-time storage health and optimization insights
System Online
GD
12%
Total Storage
23.45 TB
of 28.00 TB capacity
8%
Total Files
55
Across all tiers
Avg Health Score
63%
Needs attention
Active Alerts
3
1 critical, 1 high
Active Alerts 3
CRITICAL 2/13/2026, 4:00:00 PM
Drive WD Gold 12TB Enterprise: 28% health - Predicted failure in 5 days
Action: URGENT: Migrate all data immediately. Schedule drive replacement.
HIGH 2/13/2026, 4:00:00 PM
Drive SanDisk Extreme 1TB: 42% health - Degradation detected
Action: Schedule migration within 14 days. Enable cloud backup.
MEDIUM 2/13/2026, 4:00:00 PM
Drive WD Blue 4TB HDD: 68% health - Monitor closely
Action: Review tiering plan. Consider migrating hot data to healthier drives.

B. AI/ML Engine

- A Python-based backend service (FastAPI) that runs the XGBoost inference for drive health and K-means clustering on file metadata.

C. Optimization Logic

- Automated logic to classify Hot/Cold data.

- Cost-benefit calculator for compression and cloud migration.

D. Documentation

- Architecture Diagram: System design and data flow.

- API Documentation: Interactive Swagger/OpenAPI docs.

- Setup Guide: Comprehensive README.md for local deployment.

6. Technology Stack

- Frontend: React 19, TypeScript, Tailwind CSS, Recharts, Lucide Icons.

- Backend: Python 3.11, FastAPI, Uvicorn, Pydantic.

- ML Libraries: XGBoost, Scikit-learn, Pandas, NumPy.

- Data: SQLite (Metadata), Backblaze Metrics (Training Data), Cloud Pricing JSONs.

- DevOps: GitHub Actions (CI/CD), GitHub Pages (Hosting).

7. Future Roadmap

1. Real-Time Hardware Integration: direct interfacing with smartctl / OS kernel for live telemetry.

2. Federated Learning: Allowing the model to learn from user-specific drive degradation patterns without uploading sensitive data.

3. Ransomware Detection: Analyzing "write velocity" anomalies to detect encryption attacks in real-time.

4. Desktop Agent: An Electron-based background service for continuous local monitoring.

8. Team Details

- Team Lead: Divyanshu Patel (divyanshu.patel2023@vitstudent.ac.in)
- Member 2: Varshith Pilli
- Member 3: Ashutosh Gunjal
- Member 4: Waqar Azim
- Member 5: Soumil Gandhi